

The Alan Turing Institute

MPI Part II: Practical MPI

Baskerville Training 2025
29-30 January 2025 – David Llewellyn-Jones



Message Passing Interface

1. MPI: Message Passing Interface
2. Provides a standard for *distributed memory parallelisation*
3. Gavin already explained the theory and practice
4. Now let's apply it to some machine learning

Motivation

GPUs are great for accelerated training and inference, but processing is bound by:

1. Available memory
2. Speed of computation

How to make AI better?

1. Improved algorithms
2. Larger models
3. More training data

The last two require *more compute*

Scaling up

1. Strong motivation for scaling across GPUs
2. Single device: 4 or 8 GPUs maximum
3. Eventually want to scale across nodes
4. Get this right... the sky's the limit

Preparation

1. Open `https://docs.baskerville.ac.uk/`
2. Select Baserville Portal
3. Login if necessary
4. Select JupyterLab from the Interactive Apps list
5. Configure a 2 hour session with 1 GPU on the project `vjgo8416-training25` and the `turing` queue
6. Launch

JupyterLab - Baskerville OnDen

+

← → ↺ 🏠

🔒 📄 https://portal.baskerville.ac.uk/pun/sys/dashboard/batch_connect/sys/bc_bask_jupyter/session_contexts/new

☆ ⚙️ ⬇️ 🔄 📄 ☰

Interactive Apps

JupyterLab

GUIs

CST Studio Suite

Fiji

Linaro-Forge

RELION

cisTEM

Servers

TensorBoard

doppio

JupyterLab

This app will launch a JupyterLab server (supporting Python and Julia kernels) on Baskerville.

Kernel to load

Python 3.11.3 (2023a / GCCcore-12.3.0)

This defines the kernel you wish to load. The extra packages for use in Python can be selected from inside Jupyter using the Lmod extension. For further information please refer to the [JupyterLab documentation](#).

☐ Show Conda Environments

Include kernels for compatible Conda environments found in: `~/ .conda/ environments.txt`. For further information please see our [Conda environment documentation](#).

Number of hours

2

Number of GPUs

1

Number of GPUs to request. For each GPU you will get 25% of a node's total cores.

Baskerville Project

vjgo8416-training25

Please select the Baskerville Project to which the job will be attached.

Queue

turing


Please select the Queue/QoS on which your job will run.

Launch


Types of model parallelism

1. DDP: Distributed Data Parallel
2. FSDP: Fully Sharded Data Parallel
3. DeepSpeed ZeRO: Zero Redundancy Optimiser


Data parallel



Model parallel FSDP



Model parallel DeepSpeed




DeepSpeed

1. DeepSpeed Stage 1: optimiser state partitioning
2. DeepSpeed Stage 2: gradient partitioning
3. DeepSpeed Stage 3: model parameter partitioning

The plan


1. No more theory
2. Simplified GPT2 nano code for one GPU
3. Andrej Karpathy:
<https://youtu.be/l8pRSuU81PU>
4. Extend to support Distributed Data Parallel
5. Extend using PyTorch Lightning



The intention

1. Not to care too much about the model implementation
2. Understand changes needed to add MPI functionality
3. For this, we'll rely heavily on `diffs`
4. Using JupyterLab and SLURM

Your most important tools



$(*, /)$

Terminology

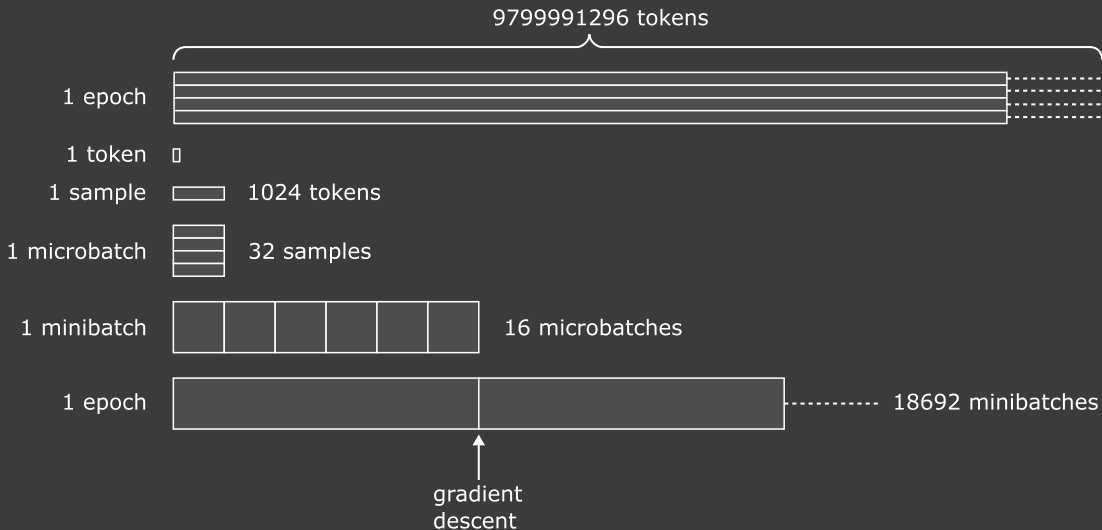
1. Vocabulary: 50257 embeddings
2. Dataset: 9799991296 FineWeb tokens, our training data
3. Sample: a sequence of 1024 tokens from the dataset
4. Microbatch: 32 samples
5. Minibatch (also called a batch): 16 microbatches, gradient accumulation performed afterwards
6. Step: the process for training on one minibatch
7. Epoch: one training sweep of the entire dataset, 18692 steps

$$18692 \times 16 \times 32 \times 1024 = 9799991296$$

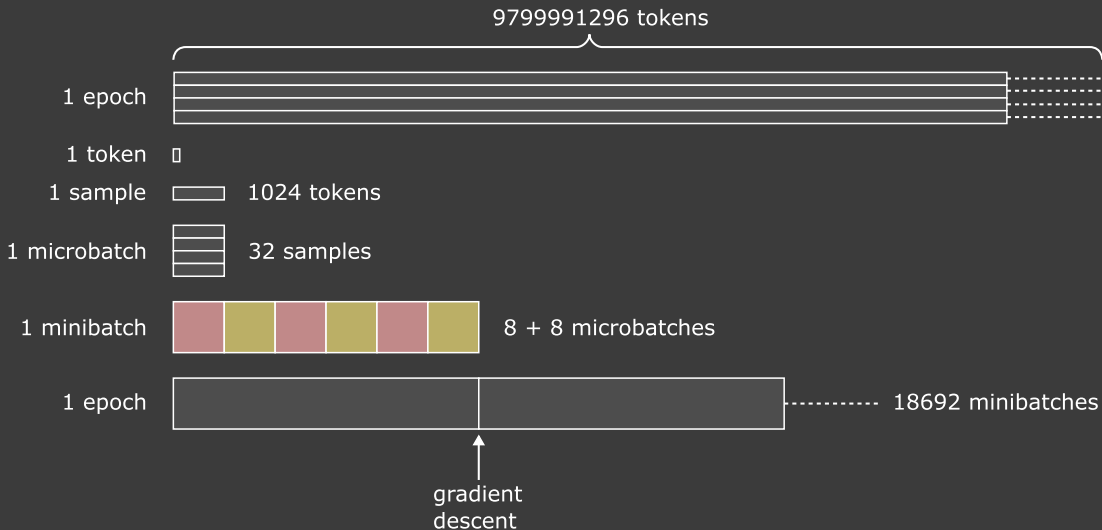
Distributed training

1. Training on n nodes, each with m GPUs
2. Total $n \times m$ GPUs
3. World size: $w = m \times n$
4. Global rank r_G is a unique index $r_G \in \{0, \dots, w - 1\}$
5. Local rank r_L is unique per device $r_L \in \{0, \dots, m - 1\}$
6. No node index, but we do have *hostnames*, e.g. bask-pg0309u05a, bask-pg0309u06a
7. We may also have multiple CPU *workers* for each node

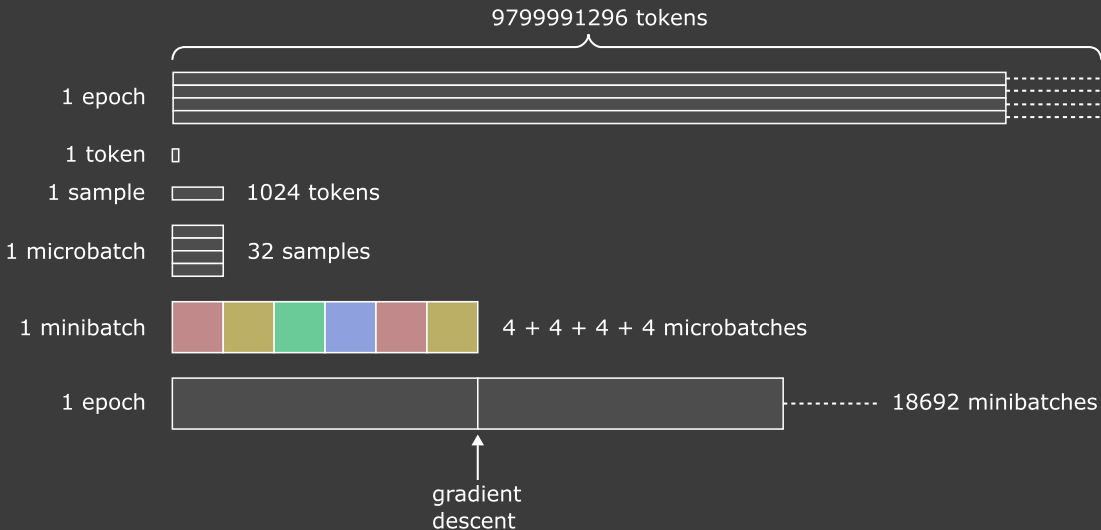
Training with 1 GPU



Training with 2 GPUs



Training with 4 GPUs



Time to look at the code

1. Open your JupyterLab page
2. Connect to Jupyter
3. Clone the repository
4. Open the `train_gpt2.py` file

```
1 # Move into your user directory
2 cd /bask/projects/v/vjgo8416-training25/$USER/
3
4 # Clone the repository
5 git clone \
6     https://github.com/alan-turing-institute/practical-mpi.git
```

JupyterLab (955888)

Host: [_bask-pg0309u06a.cluster.baskerville.ac.uk](https://>_bask-pg0309u06a.cluster.baskerville.ac.uk)

Created at: 2025-01-27 18:14:28 GMT

Time Remaining: 1 hour and 57 minutes

Session ID: [da379c88-afe6-47e7-b684-1718c3](#)

 Connect to Jupyter

My Interactive Sessions - Baskerville | train_gpt2.py - JupyterLab

https://portal.baskerville.ac.uk/node/bask-pg0309u06a.cluster.baskerville.ac.uk/47627/lab/tree/vjgo8416-karpathy/ovau2564/minGPT/gpt-2-video/train_gpt2.py

File Edit View Run Kernel Git Tabs Settings Help

Launcher | train_gpt2.py

Filter files by name

/ ... / minGPT / gpt-2-video /

| Name | Modified |
|-----------------------|-----------|
| batch-1n2g.sh | 17h ago |
| batch-ddp-1n1g... | 17h ago |
| batch-ddp-1n2g... | 17h ago |
| batch-ddp-2n2g... | 17h ago |
| batch-fineweb.sh | 17h ago |
| batch-lit-1n1g.sh | 17h ago |
| batch-lit-1n2g.sh | 17h ago |
| batch-lit-2n2g.sh | 17h ago |
| batch-lit-2n8g.sh | 17h ago |
| batch-lit-jupyter... | 18h ago |
| Calculated | yesterday |
| countfineweb.py | 5d ago |
| diff.ipynb | 17h ago |
| fineweb.py | 11d ago |
| gpt2-lit-2n2g-9... | 17h ago |
| gpt2-lit-2n2g-9... | 17h ago |
| gpu-955499-0.txt | 17h ago |
| gpu-955499-1.txt | 17h ago |
| input.txt | 11d ago |
| requirements.txt | 3d ago |
| splitfineweb.py | 17h ago |
| train_ddp.ipynb | 3d ago |
| train_gpt2_ddp.py | 17h ago |
| train_gpt2_light... | 17h ago |
| train_gpt2_test.py | 17h ago |
| train_gpt2.py | 17h ago |
| train_lightning.ip... | yesterday |
| train.ipynb | 19h ago |

```
1 #!python3
2 # vim: et:ts=4:sts=4:sw=4
3
4 from dataclasses import dataclass
5 import torch
6 import torch.nn as nn
7 from torch.nn import functional as F
8 import math
9 import tiktoken
10 import inspect
11 import os
12 import numpy as np
13
14 # -----
15
16 class CausalSelfAttention(nn.Module):
17
18     def __init__(self, config):
19         super().__init__()
20         assert config.n_embd % config.n_head == 0
21         # Key, query, value projections for all heads, but in a batch
22         self.c_attn = nn.Linear(config.n_embd, 3 * config.n_embd)
23         # Output projection
24         self.c_proj = nn.Linear(config.n_embd, config.n_embd)
25         self.c_proj.NANO_GPT_SCALE_INIT = 1
26         # Regularisation
27         self.n_head = config.n_head
28         self.n_embd = config.n_embd
29
30     def forward(self, x):
31         # Batch size, sequence length, embedding dimensionality (n_embd)
32         B, T, C = x.size()
33         # Calculate query, key, values for all heads in batch and move head forward to be the batch size
34         # nh is "number of heads", hs is "head size", and C (number of channels) = nh * hs
35         # e.g. in GPT-2 (124M), n_head=12, hs=64, so nh * hs = C = 768 channels in the Transformer
36         qkv = self.c_attn(x)
37         q, k, v = qkv.split(self.n_embd, dim=2)
```

GPT2 nano classes and functions

1. CausalSelfAttention, MLP, Block: model components
2. GPTConfig: model configuration
3. GPT: the model
4. generate(): inference
5. configure_optimizers(): training configuration
6. training_step(): one training step
7. load_tokens(): load a single FineWeb shard
8. get_shards(): find the FineWeb shards on disk
9. DataIterator: our dataset and data loader
10. get_lr(): calculate the learning rate

GPT2 nano execution

1. Set hyperparameters
2. Create model
3. Perform training loop

Training time!

1. Right click on the `train_gpt2.py` tab
2. Select **New View for Python File**
3. Open a **GPU Resources** pane from the GPU Dashboard
4. Drag the pane to the tab space on the right
5. Open a terminal in the left tab space
6. `source activate.sh`
7. `python train_gpt2.py`

My Interactive Sessions - Baskerville X train_gpt2.py - JupyterLab

https://portal.baskerville.ac.uk/node/bask-pg0309u06a.cluster.baskerville.ac.uk/47627/lab/tree/vjgo8416-karpathy/ovau2564/minGPT/gpt-2-video/train_gpt2.py

File Edit View Run Kernel Git Tabs Settings Help

Filter files by name

| Name | Modified |
|-----------------------|-----------|
| batch-ddp-1nfg... | 18h ago |
| batch-ddp-1n2g... | 18h ago |
| batch-ddp-2n2g... | 18h ago |
| batch-fineweb.sh | 18h ago |
| batch-lit-1nfg.sh | 18h ago |
| batch-lit-1n2g.sh | 18h ago |
| batch-lit-2n2g.sh | 18h ago |
| batch-lit-2n8g.sh | 18h ago |
| batch-lit-jupyter... | 18h ago |
| Calculated | yesterday |
| countfineweb.py | 5d ago |
| diff.ipynb | 18h ago |
| fineweb.py | 11d ago |
| gpt2-lit-2n2g-9... | 17h ago |
| gpt2-lit-2n2g-9... | 17h ago |
| gpu-955499-0.txt | 17h ago |
| gpu-955499-1.txt | 17h ago |
| input.txt | 11d ago |
| requirements.txt | 3d ago |
| splitfineweb.py | 18h ago |
| train_ddp.ipynb | 3d ago |
| train_gpt2_ddp.py | 17h ago |
| train_gpt2_light... | 17h ago |
| train_gpt2_test.py | 17h ago |
| train_gpt2-Copy... | 2m ago |
| train_gpt2-Copy... | 1m ago |
| train_gpt2.py | 17h ago |
| train_lightning.ip... | yesterday |

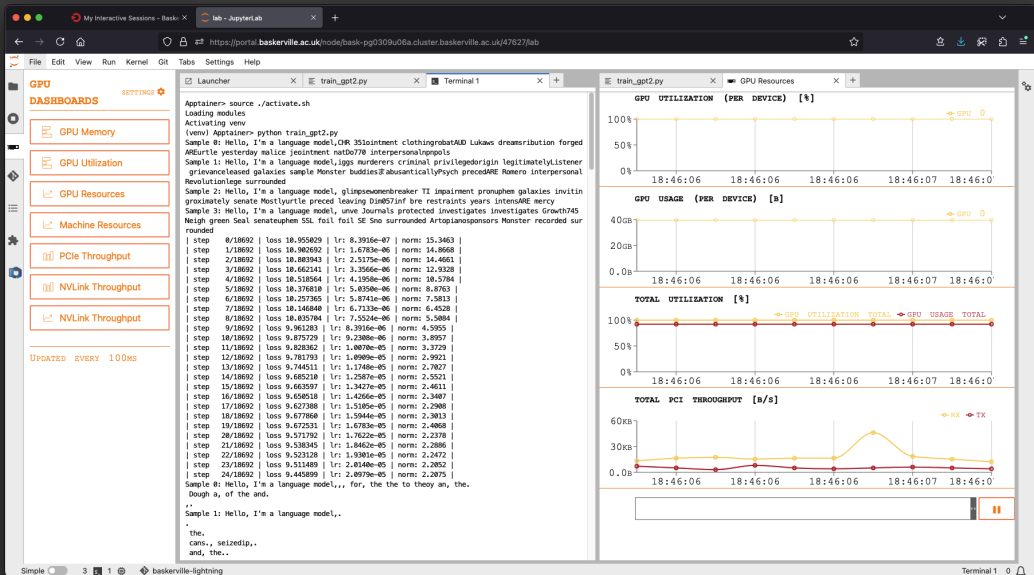
Launcher

```
263 grad_accum_steps = minibatch_size
266 max_lr = 6e-4
267 min_lr = max_lr * 0.1
268 warmup_steps = 715
269 max_steps = total_tokens // mini
270
271 assert minibatch_size % (B * T) == 0, "
microbatch size"
272
273 config = GPTConfig(vocab_size=50000,
274 train_dataset = DataIterator(B=B, T=T),
275 train_loader = iter(train_dataset_loader),
276
277 torch.set_float32_matmul_precision('high')
278
279 model = GPT(config)
280
281 model.to(device)
282
283 optimizer = model.configure_optimizers()
284
285 for step in range(max_steps):
286     if step % 25 == 0:
287         model.eval()
288         model.generate()
289
290     model.train()
291     optimizer.zero_grad()
292     loss_accum = 0.0
293     for micro_step in range(grad_accum_steps):
294         batch = next(train_loader)
295         loss = model.training_step(batch, (step * grad_accum_steps) + micro_step)
296         loss_accum += loss.detach()
297         loss.backward()
298     norm = torch.nn.utils.clip_grad_norm_(model.parameters(), 1.0)
299     lr = get_lr(step)
300     for param_group in optimizer.param_groups:
301         param_group["lr"] = lr
302     optimizer.step()
303     torch.cuda.synchronize()
304     print(f"step {step:4d}/{max_steps} | loss {loss_accum.item():0.6f} | lr: {lr:0.4e} | norm: {norm:0.4f} |")
305
306 model.eval()
307
```

Close Tab
Close All Other Tabs
Close All Tabs
Close Tabs to Right
Reopen vjgo8416-karpathy/ovau2564/minGPT/gpt-2-video/train_gpt2_test.py
Create Console for Editor
Rename Python File...
Duplicate Python File
Delete Python File
New View for Python File
Show in File Browser
Shift+Right Click for Browser Menu

```
16
17 def __init__(self, config):
18     super().__init__()
19     assert config.n_embd % config.n_head == 0
20     # Key, query, value projections for all heads, but in a batch
21     self.c_attn = nn.Linear(config.n_embd, 3 * config.n_embd)
22     # Output projection
23     self.c_proj = nn.Linear(config.n_embd, config.n_embd)
24     self.c_proj.NANO_GPT_SCALE_INIT = 1
25     # Regularisation
26     self.n_head = config.n_head
27     self.n_embd = config.n_embd
28
29 def forward(self, x):
30     # Batch size, sequence length, embedding dimensionality (n_embd)
31     B, T, C = x.size()
32     # Calculate query, key, values for all heads in batch and move head forward
33     # to the batch size
34     # nh is "number of heads", hs is "head" size, and C (number of channels) =
35     # nh * hs
36     # e.g. in GPT-2 (124M), n_head=12, hs=64, so nh * hs = C = 768 channels in
37     # the Transformer
38     qkv = self.c_attn(x)
39     q, k, v = qkv.split(self.n_embd, dim=2)
40     # (B, nh, T, hs)
41     k = k.view(B, T, self.n_head, C // self.n_head).transpose(1, 2)
42     # (B, nh, T, hs)
43     q = q.view(B, T, self.n_head, C // self.n_head).transpose(1, 2)
44     # (B, nh, T, hs)
45     v = v.view(B, T, self.n_head, C // self.n_head).transpose(1, 2)
46
```

Simple 3 Python baskerville-lightning Ln 278, Col 1 Spaces: 4 train_gpt2.py 0



Upgrade to DDP

1. Open `train_gpt2.py` in the left hand tab space
2. Open `diff_ddp.ipynb` in the right hand tab space
3. Execute the first cell of `diff_ddp.ipynb`

My Interactive Sessions - Baskerville: X diff_ddp.ipynb (3) - JupyterLab X +

https://portal.baskerville.ac.uk/node/bask-pg0309u06a.cluster.baskerville.ac.uk/47627/lab/tree/vjgo8416-karpathy/ovau2584/minGPT/gpt-2-video/diff_ddp.ipynb

File Edit View Run Kernel Git Tabs Settings Help

Launcher X train_gpt2.py X Terminal 1 X +

Filter files by name

Name Modified

- batch-ddp-1n1g... 18h ago
- batch-ddp-1n2g... 18h ago
- batch-ddp-2n2g... 18h ago
- batch-fineweb.sh 18h ago
- batch-lit-1n1g.sh 18h ago
- batch-lit-1n2g.sh 18h ago
- batch-lit-2n2g.sh 18h ago
- batch-lit-2n8g.sh 18h ago
- batch-lit-jupyter... 18h ago
- Calculated yesterday
- countfineweb.py 5d ago
- diff_ddp.ipynb 3m ago
- diff_lit.ipynb 1m ago
- fineweb.py 11d ago
- gpt2-lit-2n2g-9... 18h ago
- gpt2-lit-2n2g-9... 18h ago
- gpu-955499-0.txt 18h ago
- gpu-955499-1.txt 18h ago
- input.txt 11d ago
- requirements.txt 3d ago
- splitfineweb.py 18h ago
- train_ddp.ipynb 3d ago
- train_gpt2_ddp.py 18h ago
- train_gpt2_light... 18h ago
- train_gpt2_test.py 18h ago
- train_gpt2-Copy... 15m ago
- train_gpt2-Copy... 15m ago
- train_gpt2.py 18h ago

```
1 #!python3
2 # vim: et:ts=4:sts=4:sw=4
3
4 from dataclasses import dataclass
5 import torch
6 import torch.nn as nn
7 from torch.nn import functional as F
8 import math
9 import tiktoken
10 import inspect
11 import os
12 import numpy as np
13
14
15 class CausalSelfAttention(nn.Module):
16
17     def __init__(self, config):
18         super().__init__()
19         assert config.n_embd % config.n_head == 0
20         # Key, query, value projections for all heads, but in a batch
21         self.c_attn = nn.Linear(config.n_embd, 3 * config.n_embd)
22         # Output projection
23         self.c_proj = nn.Linear(config.n_embd, config.n_embd)
24         self.c_proj.NANO_GPT_SCALE_INIT = 1
25         # Regularisation
26         self.n_head = config.n_head
27         self.n_embd = config.n_embd
28
29     def forward(self, x):
30         # Batch size, sequence length, embedding dimensionality (n_embd)
31         B, T, C = x.size()
32         # Calculate query, key, values for all heads in batch and move head forward
33         # to be the batch size
34         # nh is "number of heads", hs is "head size", and C (number of channels) =
35         # nh * hs
36         # e.g. in GPT-2 (124M), n_head=12, hs=64, so nh * hs = C = 768 channels in
37         # the Transformer
38         qkv = self.c_attn(x)
39         q, k, v = qkv.split(self.n_embd, dim=2)
40         # (B, nh, T, hs)
41         k = k.view(B, T, self.n_head, C // self.n_head).transpose(1, 2)
42         # (B, nh, T, hs)
43         q = q.view(B, T, self.n_head, C // self.n_head).transpose(1, 2)
44         # (B, nh, T, hs)
45         v = v.view(B, T, self.n_head, C // self.n_head).transpose(1, 2)
```

train_gpt2.py GPU Resources X diff_ddp.ipynb X diff_lit.ipynb X +

Code v git Python

```
[2]: !git diff --no-index train_gpt2_test.py train_gpt2_lightning.py
diff --git a/train_gpt2_test.py b/train_gpt2_lightning.py
index 0d93334..da58190 100755
--- a/train_gpt2_test.py
+++ b/train_gpt2_lightning.py
@@ -1,9 +1,9 @@
#!python3
# vim: et:ts=4:sts=4:sw=4

-from torch.distributed import init_process_group, destroy_process_group
-from torch.nn.parallel import DistributedDataParallel as DDP
+import torch.distributed as dist
+import lightning as L
+from torch.utils.data import IterableDataset, Dataset, DataLoader
+from lightning import Trainer

from dataclasses import dataclass
import torch
@@ -97,7 +97,7 @@ class GPTConfig:
    # Embedding dimension
    n_embd: int = 768

-class GPT(nn.Module):
+class GPT(L.LightningModule):
    def __init__(self, config):
        super().__init__()
@@ -155,8 +155,8 @@ class GPT(nn.Module):
    tokens = enc.encode("Hello, I'm a language model,")
    tokens = torch.tensor(tokens, dtype=torch.long)
    tokens = tokens.unsqueeze(0).repeat(num_return_sequences, 1)
    xgen = tokens.to(device)
    sample_rng = torch.Generator(device=device)
    xgen = self.transfer_batch_to_device(tokens, self.device, 0)
    sample_rng = torch.Generator(device=self.device)
    sample_rng.manual_seed(42 + rank)
    while xgen.size(1) < max_length:
        with torch.no_grad():
@@ -185,16 +185,20 @@ class GPT(nn.Module):
    use_fused = "fused" in inspect.signature(torch.optim.AdamW).parameters
    # AdamW docs: https://pytorch.org/docs/stable/generated/torch.optim.AdamW.html
    optimizer = torch.optim.AdamW(
        train_gpt.parameters(), lr=max_lr, betas=
        (0.9, 0.95), eps=1e-8, fused=use_fused)
```

Simple 3 2 baskerville-lightning Python | Idle Mode: Command Ln 1 Col 1 diff_ddp.ipynb 0

Understanding unified diffs

1. @@ prefix indicates line numbers
@@ -before,len +after,len @@
2. + prefix in green indicates lines added
3. - prefix in red indicates lines removed
4. Replay the cell to update the diff after making changes

```
@@ -145,7 +149,7 @@ class GPT(nn.Module):
    loss = F.cross_entropy(logits.view(
s.view(-1))
    return logits, loss

- def generate(self):
+ def generate(self, rank):
    num_return_sequences = 4
    max_length = 32
    tokens = enc.encode("Hello, I'm a lang
@@ -153,7 +157,7 @@ class GPT(nn.Module):
    tokens = tokens.unsqueeze(0).repeat(nu
xgen = tokens.to(device)
    sample_rng = torch.Generator(device=de
- sample_rng.manual_seed(42)
+ sample_rng.manual_seed(42 + rank)
    while xgen.size(1) < max_length:
        with torch.no_grad():
            with torch.autocast(device_type

6):
@@ -207,22 +211,24 @@ def get_shards(split):

    class DataIterator:

- def __init__(self, B, T):
+ def __init__(self, B, T, num_processes, pr
```

Manually apply the diff – part I

1. Imports
2. One process per GPU
3. We *generate* on every process
4. One `DataIterator` per GPU
5. We must shard the data appropriately
6. Initialise the process group
7. Harvest **world size**, **rank** and **local rank** from the environment
8. Where do these come from?

Manually apply the diff – part II

9. Fix the gradient accumulation steps
10. Batch sizes must align
11. Create a DDP model
12. Use the right model variable at the right time
13. Trigger backward gradient sync
14. Perform all reduce
15. Destroy the process group

Observations

1. Most of this is boilerplate
2. The hardest part is sharding the data correctly
3. Because this is *data parallel*
4. MPI is also hard

...but it's done for us by `torch.distributed` and DDP

Training time!

1. See batch-ddp-1n1g.sh for single GPU
2. See batch-ddp-2n3g.sh for dual-node dual GPU

```
1 # Single GPU
2 python -m torch.distributed.launch \
3     --standalone --nproc_per_node=1 \
4     train_gpt2.py
5
6 # Multi-node
7 python -m torch.distributed.launch \
8     --nproc_per_node=${SLURM_GPUS_PER_NODE} \
9     --nnodes=${SLURM_NNODES} \
10    --master-port=${MASTER_PORT} --master-addr=${MASTER_ADDR} \
11    train_gpt2.py
```

Upgrade to Lightning

1. Lightning is conceptually different
2. Code is organised in `LightningModule`: init, train step, validation step, test step, optimisers
3. Code outside `LightningModule` is automated by `Trainer`
4. Remove code moving data to the GPU

Upgrade to Lightning

1. Open `train_gpt2.py` in the left hand tab space
2. Open `diff_lit.ipynb` in the right hand tab space
3. Execute the first cell of `diff_lit.ipynb`

Manually apply the diff – part I

1. Imports
2. Switch from `Module` to `LightningModule`
3. Fix the `generate()` function
4. Use the built-in `OneCycleLR` learning rate scheduler
5. Remove CUDA code from `training_step()`
6. Generate examples periodically
7. Initialise `DataIterator` using a `worker_init_fn()`
8. Remove our custom learning rate scheduler code

Manually apply the diff – part II

9. Remove the process group code
10. Lose our DDP configuration
11. Simplify the random seeding
12. Fix the steps: minibatches are now implicit
13. Use a `DataLoader` to sequence data loading
14. Remove our training loop
15. Add the `Trainer` code

Observations


1. The training code is greatly simplified
2. Lightning is SLURM-aware
3. So MPI is even easier
4. We can now switch to other strategies (*in theory*)

Training time!


1. See `batch-lit-2n8g.sh` for dual-node eight GPU
2. See `batch-lit-jupyter.sh` to queue from Jupyter

```
1 # From inside a SLURM batch script
2 srun python train_gpt2.py
3
4 # Queue execution from inside JupyterLab
5 sbatch batch-lit-jupyter.sh
```

Scaling across nodes on A100 (40 GiB) GPUs



Scaling across nodes on MX1550 (128 GiB) GPUs



Wrapping up

1. Converting code for distributed training is doable
2. Review and try out the batch scripts
3. More examples in the `hpc-landscapes` repository

<https://github.com/alan-turing-institute/hpc-landscape>